

# Commons Bonds

## Chapter 5: The Accountability Gap

The pattern repeats with such mechanical precision that once you learn to see it, you cannot stop seeing it. A company or institution extracts value from a resource, a community, or a system. The costs generated by that extraction are severed from the entity that benefited and transferred to someone else — usually someone with less power, less mobility, less ability to refuse the transfer. And when the reckoning comes, when the true scale of the severed costs becomes impossible to ignore, the original beneficiary is somehow no longer available to pay.

The reckoning always comes. The only question is who carries it.

This chapter traces that pattern across six cases separated by geography, industry, and decades, but united by the same underlying mechanism. A vermiculite mine in Montana. An oil platform in the Gulf of Mexico. A rare earth processing facility in Inner Mongolia. A national financial system's compressed double-risk-transfer in 2008. The retirement promise that was supposed to be America's most sacred compact between generations — the Social Security system itself. And the healthcare system that now routes its shortfalls through crowdfunding campaigns and unpaid caregiving labor during the most vulnerable stretches of life.

Each case tells the same story. The arithmetic is different. The mechanism is identical.

---

### Libby, Montana

The W.R. Grace Company operated a vermiculite mine in Libby, Montana, for twenty-seven years, from 1963 to 1990. The mine produced roughly eighty percent of the world's vermiculite, a mineral used for insulation and soil conditioning. The operation was profitable, generating an estimated one hundred million dollars in lifetime revenue. Grace marketed the product aggressively. Zonolite insulation filled the walls of an estimated thirty-five million American homes.

What Grace did not advertise — what Grace actively concealed — was that the vermiculite in Libby was contaminated with a particularly lethal form of asbestos.

The company knew. Internal documents released during later litigation showed that Grace had identified the contamination as early as 1963 and had measured asbestos levels in its workers throughout the mine's operation. The company knew that its workers were developing lung disease at rates that defied statistical explanation. It knew that the families of workers, exposed to dust carried home on clothing, were developing mesothelioma and asbestosis. It knew that children who played near the mine's waste piles were inhaling fibers that would kill them decades later.

And it mined for twenty-seven more years.

The human toll is still being counted. Over four hundred residents of Libby have died from asbestos-related disease. An estimated three thousand have been diagnosed with lung abnormalities consistent with asbestos exposure. The mortality rate from asbestosis in Libby was forty to eighty times the national average. The Environmental Protection Agency declared the entire town a public health emergency — the first time it had ever made such a declaration for a man-made disaster.

The financial reckoning tells the story of cost severance in numbers that cannot be argued with. Grace's mine generated roughly one hundred million dollars in revenue over its lifetime. The documented costs — medical care for the sick and dying, environmental remediation of the contaminated town, legal settlements, lost property values, and ongoing health monitoring — exceed four billion dollars. The cost-to-revenue ratio is forty to one, minimum. For every dollar Grace extracted, Libby and its residents absorbed forty dollars in costs that Grace never paid.

W.R. Grace declared bankruptcy in 2001, eleven years after the mine closed. The company's assets were reorganized. Its shareholders and executives moved on. Libby is still there, still dying, still counting its dead.

The legal architecture's response to Libby is itself a case study in cost-severance accounting. In 2005, the United States Department of Justice obtained a federal grand jury indictment of W.R. Grace and seven of its senior executives on charges of conspiracy to violate the Clean Air Act and to obstruct federal investigation of the contamination. The indictment was the largest environmental criminal prosecution in American history at the time, and it named specific executives by name and by role. In 2009, after a thirteen-week trial, all of the defendants were acquitted. The acquittal does not mean the contamination did not happen, that the executives did not know about it, or that the cost severance documented above did not occur. It means the criminal-law architecture available to the prosecution was not structurally suited to the case the prosecution was trying to make. Criminal law prosecutes individual acts, identifiable knowledge of identifiable harm at identifiable moments; cost severance is a structural mechanism that operates across decades through accumulated decisions whose individual-act elements are difficult to isolate at the criminal-knowledge-and-intent threshold. The framework's accounting does not require the criminal-law standard; it operates at the structural-mechanism level the criminal law could not reach. Libby's full record is the cost severance the framework's accounting names; the criminal acquittal is evidence that the existing legal architecture was not built to recognize the mechanism the framework prices.

---

## Deepwater Horizon

On April twentieth, 2010, the Deepwater Horizon offshore drilling platform exploded in the Gulf of Mexico, killing eleven workers and releasing an estimated 4.9 million barrels of oil into Gulf waters. The spill continued for eighty-seven days while British Petroleum, the platform's operator, attempted to cap a well that had been drilled through formations that BP's own engineers had described as "nightmare" geology.

BP had extracted substantial value from the Gulf's offshore resources. In the five years leading up to the disaster, the company's Gulf operations had generated roughly sixty billion dollars in revenue. The Macondo well itself — the specific well that blew out — was projected to produce forty to sixty million barrels of oil over its lifetime, worth between three and four billion dollars at 2010 prices.

The spill severed costs in every form cost severance can take. Spatially: the oil that had generated revenue for a British company, refined in Texas refineries, sold in markets across North America, poisoned the waters where Louisiana fishermen had worked for generations. Temporally: the damage to Gulf ecosystems persisted long after BP's quarterly reports had moved on to other projects. Ecologically: dolphins, sea turtles, and fish populations absorbed toxins that would affect reproduction and survival for decades.

The economic costs alone tell the story. BP ultimately paid twenty-one billion dollars in federal and state fines and settlements — the largest environmental penalty in United States history. The company paid an additional twenty billion dollars into a compensation fund for affected businesses and individuals. The total documented cost to BP approached sixty-five billion dollars.

Even that number understates the true cost. The National Oceanic and Atmospheric Administration estimated that the spill killed 4 to 8 billion oysters and 2 to 5 trillion newly hatched fish, along with tens of thousands of birds and sea turtles that would have supported Gulf fisheries and ecosystems for years. Tourism revenue losses in affected areas exceeded seven billion dollars in the three years following the spill. Property values along affected coastlines declined by an estimated twelve billion dollars. Aggregate estimates from academic, insurance-industry, and federal sources place total economic losses at over one hundred and fifty billion dollars.

BP's sixty billion dollars in penalties against one hundred and fifty billion dollars in total costs yields a cost recovery ratio of forty percent. The company captured the full value of the oil it extracted. It paid less than half the cost of its extraction. The framework's accountability target is one hundred percent — extractors paying the full residual cost. Deepwater paid forty. The forty-percent ratio is the documented, court-validated, settlement-ratified empirical magnitude of the accountability gap for one of the largest environmental disasters in American history.

The cost-severance asymmetry also operated across borders. The Macondo well's spill reached Mexican Gulf waters, contaminating Yucatán-Peninsula coastlines, fisheries, and tourism economies. BP's settlement with the Mexican federal government in 2018 totaled approximately twenty-five million dollars — three orders of magnitude smaller than the U.S. settlements for comparable per-mile-of-coastline damage. The asymmetry was not a function of the contamination's distribution; the slick's geographic distribution was relatively even across the affected coastlines. The asymmetry was a function of the legal architecture available to each jurisdiction's claimants. American Gulf-state litigation operated under the Oil Pollution Act of 1990 + extensive state-level environmental and tort jurisprudence + a class-action infrastructure that allowed individual claimants to aggregate at scale. Mexican claimants operated under a substantially weaker legal-architecture for cross-border environmental claims against a foreign-domiciled corporation. Same spill, same operator, same physical-mechanism of damage, structurally different bond postings. The international architectural-asymmetry is not a bug of the framework's accounting — it is a feature the accounting makes visible.

And these calculations exclude entirely the costs that cannot be monetized: the generational knowledge lost when family fishing operations went under, the cultural disruption when communities that had lived from the Gulf for centuries could no longer trust their waters, the ecological relationships severed when keystone species populations collapsed. These are real costs. They were absorbed by real people. They simply never appeared on any ledger that BP was required to balance.

---

## **Baotou, Inner Mongolia**

The city of Baotou, in China's Inner Mongolia region, produces roughly sixty percent of the world's rare earth elements — the seventeen metals whose unique magnetic and catalytic properties underwrite nearly every piece of advanced electronics on Earth. The rare earths in your phone came from Baotou. The rare earths in the wind turbine that generated the electricity you used today came from Baotou. The rare

earths in the electric vehicle that represents our best current plan for decarbonizing transportation came from Baotou.

The extraction and processing of rare earths generates enormous quantities of toxic waste. For every ton of rare earth oxides produced, the process generates approximately ten tons of wastewater containing heavy metals, acids, and radioactive materials. In Baotou, this waste flows into artificial lakes that have grown over decades of industrial activity. Residents living near these tailings ponds report elevated rates of cancer, respiratory disease, and birth defects. Agricultural land within a ten-kilometer radius of the processing facilities is no longer suitable for farming.

The companies operating in Baotou — a mix of Chinese state-owned enterprises and private firms — have extracted enormous value from global demand for rare earths. China's rare earth industry generates an estimated ten to twelve billion dollars in annual revenue, with Baotou facilities accounting for roughly half of that total. The extracted value flows through global supply chains to electronics manufacturers, automotive companies, and renewable energy firms whose products depend entirely on Baotou's rare earths.

The costs have been severed in every direction. Spatially: the toxic waste that makes Silicon Valley's technology possible sits in lakes in Inner Mongolia. The engineers designing smartphones in California never see the radioactive tailings their products require. The consumers buying those products never see the contaminated villages. Temporally: the companies purchasing rare earth feedstock pay today's extraction price; the health and environmental costs will persist for decades or centuries. Socially: the value flows to shareholders and consumers in developed economies; the costs are absorbed by rural Chinese communities with minimal political representation.

The numbers are harder to pin down than in the Montana or Gulf cases because the Chinese government restricts access to health and environmental data from the region. Some peer-reviewed research has nonetheless been published despite the access constraints — environmental epidemiology studies in *Environmental Pollution*, *Ecotoxicology and Environmental Safety*, and *Science of the Total Environment* document elevated heavy-metal concentrations (lead; cadmium; thorium; uranium) in soil, water, and resident-population biomarkers within ten kilometers of the Baotou tailings ponds; specific village-level mortality and respiratory-disease registries that researchers have compiled despite the data restrictions; and longitudinal studies tracking workforce-cohort outcomes that have, where they have been allowed to publish, found elevated rates of cancer and reproductive abnormalities consistent with the chemical and radiological exposure profiles. The published evidence is partial; the unpublished evidence is partial in the opposite direction. But even conservative estimates suggest that the long-term environmental remediation costs for Baotou's rare earth legacy will exceed one hundred billion dollars. The health care costs for affected populations — if those populations received health care equivalent to what is provided in the countries that consume their rare earths — would add tens of billions more.

Against roughly five to six billion dollars in annual value extraction, the severed costs accumulate at a rate that dwarfs what the industry has ever paid or provisioned to pay. The rare earth industry globally spends less than one percent of its revenue on environmental remediation or community health programs.

The gadgets work. The turbines spin. The electric vehicles charge. The costs accumulate in Inner Mongolia, where the rest of the world does not have to see them. And the consumer at the other end of the supply chain — the engineer in California with the smartphone in their pocket; the driver in Norway with the electric vehicle on the highway; the family in Berlin with the wind-generated electricity in their home

— is not a villain in this account. The framework’s mechanism does not require villains. The consumer is a structural participant in the cost-severance architecture by virtue of consuming the products the architecture produces; the consumer’s individual moral position is irrelevant to the structural finding. What the framework names is that the architecture severs cost from the value-extracting party; the consumer shares in the extracted value (via the product’s utility) and the consumer is structurally separated from the cost-bearing party (the Inner Mongolian families bearing the contamination). Spatial cost severance, operating through global supply chains, is what makes the consumer-and-cost-bearer relationship invisible in ways that allow the architecture to keep operating regardless of any individual consumer’s preferences.

---

## **2008 and the Double Risk Transfer**

The same mechanism operates at national-financial scale — which is where its scale becomes hardest to ignore, because unlike Appalachia or the Niger Delta, it happens to the entire economy simultaneously, leaving no cultural or geographic explanation intact.

The setup began in 1978, when Congress created the 401(k) retirement-account structure. Over the following two decades, defined-benefit pensions were systematically replaced by 401(k)s, and the replacement was presented to workers as freedom — as now owning their own retirement rather than being bound to employer-managed plans. What was not disclosed is that the freedom was primarily the corporation’s freedom from liability. Market risk, longevity risk, and behavioral risk moved from institutional balance sheets to individual workers who had no capacity to pool, hedge, or insure against them. The first transfer ran quietly for thirty years.

The second transfer came in a compressed window between 2001 and 2008, when financial institutions created mortgage-backed securities and collateralized debt obligations — instruments whose complexity was not incidental but structural, because only complex instruments could be sold to investors who could not evaluate them. The instruments were packaged into mutual funds and target-date funds sold inside the same 401(k) accounts that had absorbed the first risk transfer. When the instruments collapsed in 2008, the losses flowed into those retirement accounts. The workers who had been told in the 1980s that they were now responsible for their own retirement security bore the losses of instruments the financial sector had created, collected fees from, and would shortly be rescued from.

The rescue did come. TARP provided seven hundred billion dollars in public funds to stabilize financial institutions. The Federal Reserve’s emergency lending facilities, the full scale of which emerged only through later congressional requests and litigation, exceeded sixteen trillion dollars in cumulative emergency loans across the crisis (per the GAO’s 2011 audit). Executive bonuses at rescued institutions resumed within eighteen months. Meanwhile, approximately ten million American households received foreclosure filings between 2008 and 2012, with roughly five million resulting in completed foreclosure — more households affected than there are in the entire state of Pennsylvania, in under five years. No comparable federal intervention was organized for those households. The Home Affordable Modification Program reached a small fraction of eligible applicants and suffered administrative failures documented in multiple Government Accountability Office reports.

The asymmetry is the load-bearing fact. A system that could mobilize trillions of dollars in days to rescue institutions could not mobilize comparable resources to rescue the households whose retirement accounts had absorbed those institutions’ losses and whose homes were foreclosed in the subsequent con-

traction. The risk was transferred twice on the same money — once from corporations to workers through the pension-to-401(k) shift, and again from financial institutions to those same workers' accounts through the mortgage-backed-securities machinery. The rescue restored one side of the ledger and left the other carrying its losses. The Financial Crisis Inquiry Commission, established by Congress to investigate, concluded in 2011 that the crisis was the result of human action and inaction — not forces beyond human control. The people who made those decisions received bonuses. The people whose retirements and homes absorbed the losses received nothing.

A careful reader will object that borrowers took loans they could not afford. The serious version is straightforward: many of the mortgages that defaulted were extended to borrowers whose income could not service the payment under any reasonable model, and the borrowers signed for them. The framework's response is not that the borrowers had no agency; it is that agency requires information, and the information architecture was structurally adversarial. Subprime origination operated under a documented playbook in which originator commissions were tied to volume rather than to default rates, in which broker-disclosed terms differed materially from contract terms, in which adjustable-rate products were marketed to borrowers whose understanding of rate-reset mechanics was, by industry-internal acknowledgment, expected to be incomplete. The Consumer Financial Protection Bureau's post-crisis enforcement record documents the pattern at scale. Borrowers signed contracts; the contracts were sold under conditions designed to maximize signing rates over loan-quality. The agency was real; the information asymmetry was the lever the system pulled. Cost severance does not require borrower-as-victim framing; in the 2008 case it operated through deliberate information asymmetry — the value-extracting party had information the cost-bearing party did not, and the information structure was structurally adversarial. In other contexts the framework names — coercion, the normalization of value extraction as accepted cost of doing business, the routine functioning of subsidy as default — the lever differs. The structural finding is asymmetric cost-bearing, regardless of which mechanism produces it.

A second objection holds that 2008 was a global crisis, not a U.S.-policy failure. The serious version names the European banking exposure, the Asian export-dependency contagion, the cross-border MBS holdings, the role of the Bank for International Settlements' Basel framework in shaping global capital reserves. All of these are real. The framework's response is that the global character of the crisis amplifies rather than dissolves the U.S.-mechanism finding. The instruments that propagated the contagion were predominantly U.S.-originated MBS and CDO products. The originator-fee architecture, the rating-agency conflict-of-interest structure (issuer-pays model), and the regulatory exemptions that permitted the products to scale were U.S. policy decisions. International transmission is what gave the crisis its global magnitude; U.S.-policy origination is what gave it its instruments. A framework explaining 2008 has to engage both, and engaging both is what makes the cost-severance finding more general — not more parochial.

A third line of pushback notes that TARP was repaid, so taxpayers did not lose money in the long term. The serious version cites the U.S. Treasury accounting that records TARP's direct outlays as having been substantially recovered through repayments and asset disposals; some accountings show a modest profit. The Government Accountability Office's TARP retrospectives and the Congressional Budget Office's lifetime-cost reviews document the recovery numbers in detail. The framework's response is that TARP is the wrong accounting unit. The cost severance the chapter prices is the unrecovered welfare loss to the ten million foreclosed households, the displaced workers whose 401(k) balances reset to lower values that compounded forward at lower bases, the regional economies that contracted around foreclosure-cascade

neighborhoods, and the political-trust commons that absorbed the asymmetric-rescue pattern as evidence of the system's actual priorities. None of these losses is recovered by TARP repayment, and none of them appears on the GAO or CBO accounts that the recovery argument cites — those accounts measure federal-balance-sheet flow, not welfare-loss to cost-bearing parties. The Federal Reserve's cumulative emergency lending — the sixteen-trillion-dollar figure that emerged through congressional request — was repaid in nominal terms, but the asymmetry the rescue established is itself a cost, borne in the form of legitimacy depletion that subsequent political economy has had to operate within. The accounting question is not “did TARP cost the federal balance sheet money?” The accounting question is “did the rescue restore the cost-bearing parties to where they would have been absent the crisis?” The answer to the second question is no, by orders of magnitude, regardless of TARP's nominal recovery.

The most reformist version of the objection points to the post-crisis Dodd-Frank Wall Street Reform and Consumer Protection Act of 2010, which substantially restructured the regulatory architecture in response to the crisis, addressing the failures the foregoing analysis identifies. The serious version is real. Dodd-Frank created the Consumer Financial Protection Bureau, established the Financial Stability Oversight Council, restricted proprietary trading through the Volcker Rule, and imposed enhanced capital and liquidity requirements on systemically important institutions. These are not nothing. The architecture changed in measurable ways; the post-2010 financial system carries accountability instruments that did not exist in 2007. The framework's response is that this is exactly correct — Dodd-Frank is a partial implementation of accountability-bond architecture, and Chapter 9 will name it as one of the entering wedges by which the broader framework gets built. But the partial implementation reveals the size of the remaining gap rather than closing it. The CFPB's enforcement record exists alongside its post-2017 institutional weakening; the Volcker Rule was substantially relaxed in 2018 and 2019; the systemic-importance designations have been contested, narrowed, and in some cases reversed. The architecture moves; the architecture's stability over decades — the timeline at which intergenerational cost severance accumulates — is not yet established. A reform that holds for ten years and then is dismantled is, in cost-severance accounting, a delay rather than a correction. The framework's claim is not that nothing has been done; the framework's claim is that what has been done is small relative to the scale of the cost severance the cases of this chapter document, and that it has not been institutionally consolidated to the point where its persistence is structurally guaranteed. Reform is real; reform is partial; reform is reversible; reform that is reversible has not yet closed the accountability gap it was built to address.

What survives across each of these objections is the structural finding the chapter has been making. Borrowers had agency; the information structure was adversarial. The crisis was global; the instruments were U.S.-originated. TARP was repaid; the welfare losses on the cost-bearing side were not. Dodd-Frank reformed the architecture; the reform is partial and reversible and has not yet closed the gap.

One further observation deserves brief note before the section closes, because it bears on whether the U.S. 2008 scope-choice was inevitable or contingent. Sweden in the 1930s faced banking-crisis pressure structurally analogous to what the United States faced in 1929-1933 and again in 2008-2009, and the Swedish polity's response combined bank-rescue with substantial household-relief programs — the 1931 krona devaluation that restored export competitiveness and household-debt purchasing power; the 1933 Riksbank-coordinated bank-rescue paired with mortgage-modification programs; the 1934-1937 public-works employment programs that restored household income directly; the 1938 Saltsjöbaden Agreement framework that institutionalized the welfare-state architecture the response had built. The 1930s Swedish response is widely cited in the comparative-banking-crisis literature (Kindleberger 1973; Eichengreen 1992)

as demonstrating that bank-rescue and household-relief are not structurally either-or — a society can choose to do both. The U.S. polity in 2008-2009 and again in 2020 chose otherwise. The choice is contingent on specific historical decisions; Chapter 9 (Pricing Honestly) develops the comparison at fuller depth, including the broader cluster of cross-country and cross-century cases that demonstrate scope-choice variation in crisis-response architecture.

The 2008 case demonstrates the structural cost-severance pattern at national-financial scale. The pattern is not specific to financial markets, and the chapter has not yet finished documenting it. The cases that follow — Social Security, where the pattern operates at intergenerational scale; the healthcare commons, where it operates at the intersection of medical institutions, family caregiving, and end-of-life decision-making — show the same mechanism in different institutional contexts. The architecture differs across cases; the asymmetry between value-extractor and cost-bearer does not. What 2008 demonstrates at the scale of millions of households and trillions of dollars, Social Security demonstrates at the scale of a hundred million workers and tens of trillions of dollars in unfunded obligations, and the healthcare commons demonstrates at the scale of every American family that will eventually face an end-of-life decision under an architecture built to extract value from the dying body. Three different institutional settings; one mechanism.

---

### **Social Security: Cost Severance in the Financial Commons**

The most devastating example of cost severance in American history is hiding in plain sight, embedded in a system that most people assume represents the opposite of extraction — the promise that each generation will care for the one that came before it.

Social Security was designed in 1935 as a pay-as-you-go system. Current workers pay into the trust fund through payroll deductions. Current retirees receive benefits funded by those contributions. For decades, the system generated surpluses: more money came in from working-age Americans than went out to retirees. Those surpluses were invested in special-issue Treasury bonds, accumulating in the Social Security Trust Fund as a financial cushion for the time when the baby boom generation reached retirement age.

The baby boom generation began reaching retirement age in 2008. By 2021, Social Security was paying out more in benefits than it was collecting in payroll taxes — the demographic transition that everyone had seen coming for forty years. The trust fund has been making up the difference by redeeming its Treasury bonds. Under current law and current demographic projections, the trust fund will be depleted in 2033.

What happens in 2033 is not a mystery. Unless Congress acts, Social Security will only be able to pay seventy-seven percent of promised benefits. The program will not disappear, but it will not deliver what it has promised to the generations who built their retirement planning around those promises.

The severed cost sits in that gap. The Social Security trust-fund surpluses accumulated from 1937 onward were held in special-issue Treasury obligations rather than invested as productive sovereign capital. The opportunity cost — the gap between what the fund holds today and what it would have held if those surpluses had been invested in a broad U.S. equity index across the same window — is approximately one hundred and eight trillion dollars. That is the compounded foregone return on eighty-eight years of working-class payroll-tax contributions, extracted from the architecture by the Treasury-IOU mechanism that the surpluses were routed through.



To understand the scale: one hundred and eight trillion dollars is roughly three times the entire current federal debt. It is nearly four times the annual gross domestic product of the United States. It represents approximately three hundred thousand dollars for every person in the country — including children who will not enter the workforce for another two decades.

This is cost severance operating in the financial commons. The generation that voted for expanded Social Security benefits throughout the 1960s, 1970s, and 1980s extracted the political value of those promises — the votes, the electoral support, the sense of having provided security for their constituents. The foregone-investment cost the Treasury-IOU routing produced was severed from the politicians who chose that routing and the voters who rewarded the benefit-expansions it underwrote, and transferred forward in time to generations who were either too young to vote or not yet born.

The mechanism is identical to what happened in Libby, the Gulf, and Baotou. Value was extracted by one group of actors. Costs were severed from those actors and transferred to a different group — one with less power to refuse the transfer. The difference is that in the Social Security case, the transfer was accomplished through time rather than space, and the affected community is the working-class payroll-tax contributors whose surpluses failed to compound, plus every future American taxpayer who will absorb the resulting funding gap.

The politicians who expanded Social Security benefits without funding them extracted the electoral value and departed — through retirement, electoral defeat, or death — before the costs came due. The voters who supported those expansions received the benefits they were promised. The foregone compounding was transferred to their children and grandchildren, who will inherit a system whose surpluses were routed through the Treasury-IOU mechanism instead of into productive sovereign capital, and face the political choice between absorbing the resulting funding gap themselves or breaking commitments that older Americans organized their lives around.

Social Security is not a Ponzi scheme — it is a genuine intergenerational transfer program providing essential benefits to millions of Americans. But the specific way it has been funded reveals cost severance operating at civilizational scale. One hundred and eight trillion dollars in foregone return, extracted from the architecture by the routing choice that earlier generations made and transferred to generations who had no say in making it. The largest cost severance in human history, hiding inside a program that Americans think of as the most sacred commitment one generation can make to another.

The architectural choice the United States made in 1935 — pay-as-you-go retirement security funded by current-worker payroll deductions, with surpluses held as Treasury obligations rather than invested as productive sovereign capital — was a contingent choice, not an inevitable one. Chapter 4 has documented Norway's contrasting architectural choice for analogous intergenerational fiscal architecture: the Government Pension Fund Global, which channels petroleum rents into a sovereign-fund structure that compounds across decades and accumulates capital on behalf of future generations rather than depleting it. The two architectures address structurally similar problems and produce structurally different outcomes; the difference is not a function of national wealth or demographic pressure but of architectural choice made at specific historical moments. The chapter on Norway documents what was possible; the present chapter documents what was, and continues to be, the cost severance the U.S. choice produces.

As with the 2008 case, the Social Security account invites two further objections worth engaging. The first is a public-finance-prudence objection: the Treasury-bond mechanism the trust fund used was administratively defensible, not the source of severance. The second is a timeline-and-tractability objection:

the funding gap is real but solvable through minor adjustments at a manageable horizon, and the framework's framing inflates what is structurally a soluble policy problem.

The public-finance-prudence objection holds that the Treasury-bond mechanism by which Social Security surpluses were accumulated was prudent, not opportunistic — that holding the trust fund's reserves in special-issue Treasury obligations is exactly what a sound sovereign-fund manager would do, that the alternative of holding the reserves in equities or other risk assets would have exposed retirees' security to market volatility, and that the framework's account treats a defensible administrative choice as if it were the mechanism of extraction.

The serious form of the objection has three parts. First, Treasury bonds are the safest asset class available to a U.S. sovereign fund — they default only if the federal government defaults, which is the same condition under which any U.S.-domiciled retirement asset would fail. Second, the General Fund borrowing that the trust fund's bond holdings facilitated did not come at Social Security's expense; the bonds bear interest at competitive rates, which is the same return the trust fund would have earned holding any other Treasury security. Third, the alternative — holding trust fund surpluses in equities or private assets — would have introduced both market risk and the political-economy risk that comes from a federal entity holding controlling interests in private firms, a pattern most contemporary public-finance economists treat as a more serious problem than the bond-holding pattern that emerged.

Each of these is sound on its own terms. The framework's response is that they are addressing a different question than the one the chapter raises. The trust-fund bond-holding pattern is administratively prudent at the level of the trust fund's portfolio. The cost severance the chapter prices is structurally upstream of the portfolio question entirely.

The cost severance is not located in how the surpluses were invested. It is located in how the obligations were generated in the first place. Throughout the nineteen-sixties, seventies, and eighties, Social Security benefit expansions were enacted by Congresses whose members would, in nearly every case, have left office before the demographic-transition obligations matured. The benefit expansions were funded by the demographic structure of the moment — the unusually large working-age cohort relative to the retirement-age cohort during exactly the decades the expansions were enacted — and the funding gap that demographic transition would eventually create was, in every actuarial projection issued during that period, openly visible. The trust-fund Treasury-bond mechanism was the administrative vehicle through which the surplus-then-deficit demographic curve was managed. It did not create the gap. The political incentive to expand benefits during a demographic-surplus window without funding the demographic-deficit window that the actuarial reports identified as inevitable created the gap. The bond-holding pattern is downstream of that political incentive; it is not the source of severance.

The framework's account is not that the bonds were imprudent. It is that the obligations were generated under a structure in which value-extraction (the electoral support generated by benefit expansions) was severed from cost-bearing (the future-generation taxpayers who would actually fund the demographic-transition deficit) by exactly the time-displacement mechanism the framework names. Holding the surpluses in Treasuries during the surplus-window was prudent. Generating obligations during the surplus-window that the actuarial reports projected would deplete during the deficit-window was the cost severance. The two questions operate at different levels, and a defense at the portfolio level does not address the obligation-generation level the chapter prices.

The timeline-and-tractability objection holds that the framework's account inflates what is, on a sober

reading, a soluble public-finance problem. The trust fund will not run out tomorrow; it will run out in 2033, more than half a century after the legislative expansions that generated the obligations. Even after depletion, payroll-tax revenue continues to fund roughly seventy-seven percent of promised benefits. The Trustees' annual reports identify a finite range of policy adjustments — payroll tax rate increases of one to two percentage points; cap removals on the taxable wage base; benefit-formula adjustments at the upper income range; eligibility-age increases tracking longevity — any combination of which closes the seventy-five-year actuarial gap with substantially less disruption than the framing suggests. The Bipartisan Policy Center, the Concord Coalition, and successive iterations of the Simpson-Bowles and Domenici-Rivlin commissions have produced specific adjustment packages. The objection holds that the gap is a tractable policy-design problem, not a civilizational-scale severance.

The serious version of the objection is correct as far as it goes. A combination of payroll-tax rate adjustments, taxable-wage-base cap modifications, benefit-formula refinements, and eligibility-age changes can mathematically close the seventy-five-year gap. The Trustees' actuarial work is robust and the policy proposals are well-developed. A society that adjusted at the timeline this objection identifies would not face a Social Security collapse; it would face a phased-in adjustment.

The framework's response is that the timeline question and the obligation-generation question operate at different levels, and that closing the timeline gap through adjustment does not address what the framework prices. The framework prices the severance — the structural pattern by which value-extracting parties (the cohorts who voted for benefit expansions and the politicians who enacted them) were systematically separated from cost-bearing parties (the future-generation taxpayers who would fund the demographic deficit). The timeline objection grants that the gap is closeable through adjustments at the cost-bearing-party's expense; that is exactly what the framework names as the structural problem. Closing the gap by raising payroll taxes on the demographic-deficit-window cohort, or by reducing benefits to retirees in that window, or by raising the eligibility age for the cohort that has been paying into the system for forty years, transfers the obligation onto the cost-bearing-party — not onto the cohorts who extracted the electoral value of generating the obligation. The gap is mathematically closeable; the severance the closure achieves is the severance the framework names. Sober public-finance adjustment is what cost severance looks like at civilizational scale when the time-displacement runs its course. The objection's solution is what the framework's account predicted would happen: when the obligations mature, the cost-bearing-party absorbs the adjustment because the value-extracting-party is no longer accessible. Tractability is not absence of severance; it is severance unfolding on schedule.

What survives across both objections is the chapter's structural finding. The trust-fund bond-holding mechanism was administratively defensible. The political-economy mechanism that generated the obligations the trust fund was holding bonds against was not. The funding gap is mathematically closeable through adjustment; the adjustments themselves transfer the obligation onto the cost-bearing-party rather than onto the value-extracting-party. The first observation does not absolve the second; the second observation does not refute the framework's account. The accountability gap is upstream of the portfolio question and downstream of the timeline question — it sits in the obligation-generation mechanism that both questions take as given.

The Social Security case demonstrates the cost-severance pattern at intergenerational scale, where time itself is the distance that separates value-extraction from cost-bearing. The 2008 case demonstrated the same pattern at national-financial scale, where institutional architecture chose what was in-scope for protection and what was not. The healthcare commons that the next section examines demonstrates the

pattern at a third register: the institutional architecture of medical financing produces cost severance on individual bodies during the most vulnerable stretches of life, and the costs it severs are not abstract externalities but the specific financial ruin of specific households. Three different institutional settings; three different scales at which the severance gets hidden — financial-system architecture; intergenerational time; institutional-medical-financing — and one underlying mechanism. The chapter's structural argument is that the mechanism is not specific to any particular institutional context; it is structural to the absence of accountability instruments that would internalize cost. The healthcare case completes the chapter's documentation; the synthesis that closes the chapter then names the pattern explicitly across all six cases.

---

## The Healthcare Commons

The mechanism operates in American healthcare with a particular cruelty, because it operates on individual bodies during the most vulnerable stretches of life, and because the costs it severs are not abstract externalities but the specific financial ruin of specific households.

The most visible indicator is a growth curve. Between 2011 and 2020, the annual number of US medical fundraising campaigns on the platform GoFundMe grew twenty-five-fold. By 2020, more than a third of all GoFundMe campaigns were for medical expenses; between 2016 and 2020, 438,000 such campaigns raised over two billion dollars from 21.7 million donations. Only twelve percent reached their funding goals. Sixteen percent received no donations at all. GoFundMe's own chief executive has acknowledged the platform was never designed to function as healthcare financing and has become the default. A safety net with an eighty-eight-percent failure rate is not a safety net. It is a documentation system for a severance the formal system has stopped trying to prevent.

The severance lands unevenly. Medical fundraising campaigns organized by individuals with distinctively African American or Hispanic surnames raise substantially less and are significantly less likely to meet their goals, even controlling for geography and time period. The replacement for a public safety net replicates the inequities the original net had been designed, however imperfectly, to reduce. This is cost severance operating through a second layer — the formal system drops the medical cost; the informal backup does not catch equally by race; the severed cost concentrates on populations whose existing wealth-and-network constraints leave them with the least capacity to absorb it.

The structural pressure on households is not solely a function of healthcare-cost growth in absolute terms. It is also a function of wage stagnation that has not kept pace with healthcare-cost growth across the same decades. Lawrence Mishel and Josh Bivens at the Economic Policy Institute have documented that the median U.S. wage tracked roughly with productivity growth from the late 1940s through the early 1970s, then decoupled — productivity continued to rise; the median wage stagnated in real terms for the next half century. Over the same half century, healthcare costs grew at multiples of general inflation. The structural pressure on households is the gap between these two curves: wages flat against productivity that wage-earners produced; healthcare costs climbing against the same wages. Jacob Hacker's *Great Risk Shift* (2006) named the broader pattern of risk-and-cost transfer from institutional to individual balance sheets that healthcare-cost-severance is one specific instance of. The GoFundMe-as-default-safety-net pattern is what cost severance looks like when the transfer mechanism's cost-bearing party is structurally unable to absorb the transfer's magnitude — wages stagnant, savings depleted, household-debt accumulating, fundraising-from-strangers as the residual mechanism.

At the other end of the arc, in the individual case the journalist Katy Butler documented, Medicare reimbursement pays a physician fifty-four dollars for a long office visit to discuss treatment alternatives, and four hundred and sixty-one dollars — on top of a roughly twelve-thousand-dollar hospital flat fee, nearly half of which goes to the device manufacturer — for the implantation of a pacemaker. The liability architecture points in the same direction. Performing the procedure protects against malpractice exposure; declining it creates exposure. Both incentives arrive at the same outcome. Butler's father spent five years in profound cognitive decline with the pacemaker still functioning, cared for by her eighty-year-old mother in eighty-hour weeks of unpaid labor for which no compensation exists and no ledger captures. The device manufacturer received approximately six thousand dollars. The conversation that might have prevented the procedure was worth fifty-four. This is cost severance operating on the dying body: profit extracted by the cardiologist, the hospital, and the device manufacturer; cost externalized onto the family as unpaid caregiving, onto the patient as prolonged decline, and onto an eighty-year-old spouse absorbing the labor that no formal architecture existed to account for.

Butler's case is one archetype of a systemic pattern. Atul Gawande's *Being Mortal* (2014) documents the broader structure: the American healthcare-financing architecture systematically over-treats dying patients (because the reimbursement architecture incentivizes intervention at every decision point) while substantially under-funding palliative-care infrastructure (because palliative-care reimbursement is a fraction of intervention reimbursement, and most American medical-school curricula carry minimal palliative-care training). The systemic outcome is a population of dying patients receiving medically-unnecessary procedures whose cost-severance pattern reproduces Butler's case at scale: the device manufacturer extracts value; the hospital extracts value; the cardiology practice extracts value; the dying patient and their family bear the cost. The architecture does not require any individual cardiologist or hospital administrator to choose extraction over compassion at any individual decision point — the architecture's reimbursement structure produces the pattern by aggregating individual rational responses to incentives that are themselves the policy-design choice the framework names. End-of-life care is the most emotionally legible illustration of cost severance in healthcare; it is not the only domain where the same architecture operates. Maternal-care reimbursement, mental-health-care reimbursement, chronic-disease-management reimbursement all reproduce the same pattern at different points along the arc of human life.

The healthcare commons demonstrates the cost-severance pattern at the institutional-medical-financing scale, where the gap between value-extractor and cost-bearer is collapsed to a single physician-patient encounter and then re-expanded through the architecture of reimbursement, malpractice liability, and family-caregiving labor that no formal compensation mechanism captures. The cost-severance the framework names is operating on individual bodies at the points in life where they are least able to refuse the transfer, and the costs it severs are not abstract externalities but the specific financial ruin and unpaid labor of specific households. Among the wealthy nations, the United States is the only one whose healthcare-financing architecture produces the GoFundMe-as-default pattern at this scale; the cross-national comparison Chapter 9 develops at fuller depth will return to this finding.

---

### The Pattern Made Visible

Six cases. Six industries. Four continents, two of them represented twice. Six entirely different types of extraction. The pattern is identical across all of them.

An entity — a company, a generation, a political system — extracts value from a resource or system that could theoretically be managed for long-term sustainability. That value extraction generates costs: environmental damage, health impacts, financial obligations, the permanent foreclosure of future options. Those costs are severed from the entity that created them and transferred to parties who had minimal power to prevent the transfer: workers, communities, ecosystems, future generations.

When the costs eventually materialize — when the workers get sick, when the environment collapses, when the bills come due — the original beneficiary is protected from the full reckoning by some combination of distance, time, asymmetry of what is known, asymmetry of who can refuse, normalization of the cost as routine, and the appearance of abundance that prevents the cost from registering at all. The costs are absorbed by those who remain: the people who cannot leave, the systems that cannot relocate, the generations who inherit the mess.

In Libby, the costs were absorbed by a Montana town that could not move away from the contamination Grace left behind. In the Gulf, the costs were absorbed by ecosystems that could not relocate and communities whose livelihoods were tied to waters they could not abandon. In Baotou, the costs are absorbed by rural Chinese populations with limited political voice and no practical ability to leave. In the 2008 crisis, the costs were absorbed by the ten million households whose homes went into foreclosure and the workers whose retirement accounts carried losses no institutional rescue reached. In Social Security, the costs will be absorbed by future Americans who cannot opt out of inheriting the fiscal obligations that previous generations created. In the healthcare commons, the costs are absorbed by households that fail to fund their GoFundMe campaigns and by eighty-year-old spouses providing caregiving labor that appears on no ledger.

The companies declare bankruptcy. The politicians leave office. The generations who made the promises pass away. The costs remain, transferred to whoever cannot escape them.

This is not a failure of the system. This is the system working exactly as it has been designed to work — not by conscious intention, but by the accumulated logic of structures that reward value extraction and externalize costs to those with the least power to resist. Cost severance is not an accident. It is not an unintended consequence. It is the mechanism by which every extractive economy has ever operated, and it is still operating today.

The question is not whether this is happening — the six cases in this chapter prove that it is. The question is what it looks like to build accountability structures that can interrupt the pattern before the costs are severed, rather than trying to recover them after the damage is done and the beneficiaries have already departed. Brett Christophers' *The Price is Wrong* (2024), engaged at depth in Chapter 9, develops the contemporary climate-finance argument that profit, not price, drives investment — a finding the next chapter returns to when it engages property-rights limits and the political economy of resistance.

### **Restitution and Foreclosure: Two Directions of Accountability**

The cases this chapter has documented share a structural pattern, and the chapter's previous synthesis named that pattern. What remains is to name the framework's response apparatus: the specific accounting instruments the framework proposes when accountability architecture is being built rather than diagnosed. The remainder of this chapter introduces those instruments at the level of structure — the methodology that follows in Chapter 6 will formalize them; the policy-economy work in Chapter 9 will engage what their political implementation looks like; the present chapter only names what the apparatus is and

what established academic-and-legal traditions it operates within. The naming matters because the cases the chapter has documented produce a question — *what would honest accounting of these severances actually require?* — and the framework’s two-instrument response is the answer that the remaining chapters elaborate.

The accountability gap names a problem with two temporal faces. There is the harm already done — the McDowell life-expectancy gap, the Libby asbestos deaths, the Gulf ecosystems that have not recovered, the foreclosed households that have not regained the homes they lost. And there is the harm not yet done — the future-generation climate cost of carbon emitted today, the petroleum hydrocarbons that will be foreclosed for non-energy uses fifty years from now, the Social Security obligation a younger cohort has not yet been asked to bear in full. An accountability architecture that addresses one direction without the other addresses only half the problem.

The framework’s response to the bidirectional structure is two coordinated instruments. The first — the Restitution Bond — addresses harm already realized. The bond’s posting discharges the historical extraction-portfolio: the cumulative damage borne by communities, workers, and ecosystems whose costs were severed at extraction-time. The second — the Foreclosure Bond — addresses harm yet to be realized. The bond’s posting discharges potential future-generation costs that current extraction is foreclosing on. Together, they form the aggregate accountability bond the framework’s central equation requires when it states that Cost Severance equals Residual Commons Value minus Bond posting.

The framework’s choice of *bond* as the umbrella term across its accountability instruments is deliberate. Other terms were available — *tax, fee, fund* — and each was rejected for a specific structural reason. Taxes are paid at the moment of transaction, against general revenue, with no architecture for matching a posted dollar amount to a specific future obligation. Fees license activity rather than discharge harm. Funds aggregate without discharging. A bond, by contrast, performs three operations the framework requires simultaneously: it is posted before the full cost has been borne, it carries forward across the long tail of the cost it is meant to cover, and it discharges — observably, by a specific dollar figure — against a specific liability when the obligation comes due. The reclamation bonds mining companies post under the Surface Mining Control and Reclamation Act do this. The decommissioning bonds oil-and-gas operators post against well closure do this. The sovereign wealth fund Norway built against its petroleum extraction does this on a national scale. The reparations economics Darity and Mullen developed treats the obligation as a documented historical debt — bond-like in structure if not yet in name. The civic-republican tradition Pettit and Skinner recovered names a standing obligation against future generations that bond posting is the natural accounting form for. The traditions converge on the instrument. The framework uses the term in continuity with these traditions, and in the meaning they have long built into it.

A reader engaging the framework’s *Restitution Bond* alongside the chapter’s lineage prose will notice an asymmetry the prose has not yet engaged: the bond is named *Restitution*, the tradition that grounds the methodology is named *Reparations*. The asymmetry is intentional and warrants the explanation. *Restitution* is the older legal-and-philosophical term for the obligation to restore what was taken — *restitutio in integrum*, restoration to the prior condition or its monetary equivalent — and the framework needs an umbrella broad enough to cover the full field of cost-severance contexts the bond will be posted against. The Coates-Darity-Mullen-Hamilton-Conley reparations tradition is the framework’s primary informing source; its methodology and rhetorical architecture are how the framework knows what it knows about backward-looking accountability. But *Reparations Bond* would narrow the umbrella to one historically-specific application and would make the methodology’s portability to environmental restitution (Libby

asbestos), to occupational-disease restitution (the black lung trust fund), to community-restitution (McDowell), and to ecosystem-restitution (the Niger Delta) read as analogical extension rather than as direct instrument-application. The narrower alternatives fail differently. *Restoration Bond* presupposes physical restorability that many cost-severance harms — lives lost, ecosystems extinguished, communities dispersed — do not allow. *Indemnification Bond* routes through a third-party insurance register the framework does not operate in. *Damages Bond* collapses the bond instrument into the measurement instrument the framework already names separately. *Restitution Bond* covers what the alternatives cannot, while leaving room for the reparations tradition to retain its specific historical claim. The framework joins the tradition rather than translating it; the bond's name marks the joining at the instrument level while the methodology, in its core moves, learns from where the tradition has already done the work.

The Restitution Bond operates within an established academic and legal lineage. The argument-architecture for reparations claims grounded in structural-economic harm, rather than tortious individual acts, was articulated most influentially by Ta-Nehisi Coates in *The Case for Reparations* (*The Atlantic*, 2014) — an essay whose form is now widely recognized as having returned reparations to mainstream American political discourse. Coates's framing pattern — name the specific mechanism by which harm was produced, name the specific harm, name the specific debt — is the rhetorical companion to what reparations economics develops methodologically. The framework's CSD (Cost Severance Damages) instrument is a structural extension of Coates's argument-architecture into extraction-community contexts: name the specific extraction mechanism (cost severance), name the specific harm (the documented life-expectancy gap, the institutional collapse, the cumulative health-care costs externalized to families and Medicare), name the specific debt (CSD as backward-looking accounting).

Reparations economics develops Coates's framing into formal methodology. William Darity Jr. and A. Kirsten Mullen's *From Here to Equality: Reparations for Black Americans in the Twenty-First Century* (2020) is the contemporary landmark — their wealth-gap calculation methodology is the careful accumulation of specific economic mechanisms by which intergenerational harm has been transferred and compounded. Hamilton, Darity, Price, Sridharan, and Tippet's *Umbrellas Don't Make It Rain* (2015) extends the methodology empirically: matched-comparison work across achievement categories demonstrating that conventional explanations for the racial wealth gap (educational attainment, employment, family structure) cannot account for the gap's magnitude — even Black households at the highest achievement strata display median wealth substantially below white households at lower achievement strata. The structural variable is wealth, not behavior. Dalton Conley's earlier *Being Black, Living in the Red* (1999) established the wealth-as-structural-variable methodology these later works develop. The framework's CSD instrument applies the same matched-comparison discipline to extraction-community contexts: McDowell County's wealth collapse against comparable non-extraction Appalachian counties; Libby's cumulative health-cost burden against comparable communities without W.R. Grace's asbestos exposure; the opioid-Appalachia mortality demography against comparable counties where Purdue Pharma's documented marketing-target mapping did not concentrate. The methodology is portable; the per-context calibration is what differs.

The legal-architecture lineage the framework engages is most rigorously developed in Katharina Pistor's *The Code of Capital* (2019). Pistor demonstrates that capital is not a natural asset class but a legal artifact: specific legal modules — property law, collateral law, trust law, corporate veil law, bankruptcy law, and choice-of-law clauses — encode assets as capital and protect that capital from accountability across jurisdictions. The Libby case is exactly the pattern she names. W.R. Grace's bankruptcy filing was not an accident



but a structural feature of the legal architecture: bankruptcy law allows the corporate-veil-protected entity to discharge liabilities by reorganizing, transferring claims onto creditors and the public. The 2009 acquittal was the criminal-law architecture failing to reach what Pistor names as a legal-architectural design feature, not a bug. The framework's accountability-bond apparatus is the Pistor-adjacent move: where she diagnoses the legal modules that encode cost severance, the framework prices what the modules sever.

A reader meeting *Foreclosure Bond* for the first time will hear the second word first, and most readers will hear it through the 2008 housing crisis — the wave of mortgage foreclosures that took primary assets out of household balance sheets across the country, with disproportionate concentration in Black and working-class communities. The framework names that association as continuity, not coincidence. Foreclosure derives from the Old French *forclore* — to bar from — and the mortgage-foreclosure event is one institutional context where the bar-from mechanism operates: a household is barred from continued access to a primary asset by a legal architecture that allows the bar-from to occur at scale. The framework names a structurally cousin instrument: future generations are barred from continued access to commons-options — habitable atmosphere, intact ecosystems, reserves of irreplaceable materials — by a legal-and-economic architecture that allows the bar-from to occur across time. Same mechanism. Different substrate. The Foreclosure Bond is the accountability instrument the framework proposes against the second bar-from. Alternative names were available and each fails for a specific reason. *Future-Generations Bond* loses the bar-from structural feature and reads as soft-philanthropic. *Sustainability Bond* has been co-opted by ESG-finance markets where the term has lost specificity. *Reclamation Bond* is already an established term in the Surface Mining Control and Reclamation Act regime the framework cites as part of the Foreclosure Bond's lineage, but reclamation names the act of restoring a site after extraction — a subset of what the Foreclosure Bond covers, not the umbrella. *Conservation Bond* preserves the resource without naming the bar-from-future-options structural feature. *Intergenerational Bond* describes the temporal direction without naming the mechanism. *Foreclosure Bond* names the mechanism the others describe; the cousin-association the reader brings to the term is the right cousin to bring.

The Foreclosure Bond operates in adjacent but distinct territory. The lineage here is environmental-bond economics — Hartwick's rule for sustainable extraction (1977), the reclamation-bond tradition that mining and oil-and-gas regulation in the United States has implemented at varying levels of seriousness for half a century, and the sovereign wealth fund architecture that Norway built for its petroleum extraction. These are not reparations instruments. They are foreclosure-prevention instruments — accountability postings made against ongoing or proposed extraction whose costs would otherwise be inherited by future generations whose moral standing the framework treats seriously (Parfit 1984's intergenerational impersonal-outcomes framework grounds this commitment philosophically).

The Foreclosure Bond also operates within a political-philosophical lineage worth naming explicitly. The civic-republican tradition — most contemporarily articulated in Philip Pettit's *Republicanism* (1997) and historically recovered in Quentin Skinner's *Liberty Before Liberalism* (1998) — frames freedom as non-domination, and non-domination as requiring institutional architecture that future generations cannot trivially be subjected to. A society that extracts now in ways that subject future generations to dominated-conditions (climate destabilized by present emissions; commons foreclosed for present consumption; political-economic architectures rigged before future generations have political voice) violates the civic-republican commitment regardless of whether any specific contract has been broken. The Foreclosure Bond is the framework's accounting instrument for that commitment: a society that does not post such a bond against ongoing extraction is, in civic-republican terms, dominating its future

generations through the architecture of present consumption. The framework does not require the reader to accept civic-republican political philosophy as the only frame; the bond instrument operates regardless of which political-philosophical tradition supplies the moral standing of future generations. But the civic-republican framing is a substantial tradition that grounds the Foreclosure Bond's core commitment, and the framework engages it explicitly under the discipline of accommodating where alternative traditions converge on the same accountability requirement.

The two instruments are structurally independent. A community whose commons has already been largely consumed — McDowell after the coal era, the Niger Delta after sixty years of extraction — carries a high Restitution Bond obligation and a relatively lower Foreclosure Bond obligation, because the foreclosure has already happened. A context where extraction is proposed but has not yet occurred — asteroid mining as an ongoing thought experiment, deep-seabed mining as an emerging policy debate — carries a low Restitution Bond obligation and a high Foreclosure Bond obligation, because the foreclosure is what is being decided about now. Most real cases carry both: Norwegian petroleum extraction internalized substantial restitution through the welfare-state architecture for Norwegian citizens but did not internalize the foreclosure cost for non-Norwegian populations affected by Norwegian oil's atmospheric externalities. Norway carries a paid Restitution Bond on its citizens; the Foreclosure Bond is globally underdeveloped. The decomposition is the apparatus that makes both visible simultaneously.

The methodology this book has been demonstrating runs in one direction. Each case study — McDowell's coal, Norway's oil, the Macondo well, Libby vermiculite, Baotou's tailings — has been priced forward. The framework's apparatus — the four gates and the three ways of counting — does that work. The Technical Appendix §5.5 develops both directions formally; Chapter 6 develops each.

The same apparatus runs in reverse.

In later chapters we will walk through a theoretical use case of the abundance of clean air being stripped — by an off-world colony on Mars engineering its own atmosphere, by a climate crisis on Earth — costs of atmospheric services become visible that the abundance had masked. The framework gates and triangulates them. This is the forward case. When the abundance of freedom is stripped — by slavery, by forced labor, by the variants of coercion that have organized large parts of human history — the costs of the autonomy commons surface that the abundance had hidden. The framework's apparatus reaches these costs too. They yield to the same four gates. They triangulate through the same three methods, applied backward: **Method 1 — Replacement Cost** (here: remediation cost rather than substitute-engineering cost); **Method 2 — Revealed Preference** (here: revealed restitution rather than revealed restraint); **Method 3 — Scarcity-Adjusted Option Value** (here: the option value extinguished at the time of past extraction, evaluable from what we know now).

The most developed backward-direction work in the contemporary literature is Darity and Mullen's *From Here to Equality* (2020), which prices the racial wealth gap accumulated through American policy choices since 1619. The methodology is the framework's, applied backward, with Darity and Mullen as the primary developers. Their more recent work prices the longevity gap — the six-to-seven year average life-expectancy differential between Black-American and white-American populations — as a measurable legacy effect. McDowell County's thirteen-year life-expectancy gap, cited in earlier chapters, is the same legacy-effect variant applied to coal-extraction-affected populations.

A note on vocabulary. Darity and Mullen structure reparations as a three-component program: acknowledgment-and-atonement; *redress*, decomposing into symbolic redress and material restitution

(they declare in favor of restitution); and closure. The framework's *Restitution Bond* is positioned, then, not as a concept alongside reparations but as the instrument that operationalizes the restitution component of redress within their framework, extended to the broader case-class this book engages. When the framework calls something a Restitution Bond, it uses their language inside their typology, not against it.

Some backward variables sit at the framework's analytical boundary rather than inside it. The coercion vector itself — what should be assigned to coerced labor as a cost dimension, distinct from coerced labor's legacy effects — remains methodologically unresolved in the reparations-economics field. Slavery's structural distinctness from the cigarette-industry case, the Libby asbestos case, the opioid-Appalachia case, is that those cases operate through information-asymmetry: workers and consumers proceeded because they did not know what they would bear. Coerced extraction is different: the cost-bearer knows, and is forced to bear it anyway. The framework reaches coerced cases through their legacy effects — wealth gaps, longevity gaps, cultural-knowledge transmission disruption, intergenerational trauma. The direct pricing of the coercion vector remains an open question, acknowledged outside the book's empirical scope; the Technical Appendix's scope-of-applicability boundary at §1.10 names the variable explicitly (informed by Darity 2026).

Indigenous land dispossession, cultural-knowledge severance through colonization, the multi-generational accumulation of cohesion-commons damage — each sits in the same territory: structurally accessible to the framework's methodology applied backward, computed in subsequent work rather than this book.

This is by design. The book demonstrates the framework forward; the backward application is structurally identical methodology, carried in formal articulation in the Technical Appendix (§5.5). Voices in the reparations-economics tradition — Darity, Mullen, Hamilton, Coates, Anderson, and the broader field — have been doing this work for years; the framework recognizes and defers to their work in identifying and pricing the variables.

What this chapter has documented is the empirical state of bond posting in the American extraction regimes: the bond posted is far below the residual commons value, in nearly every case examined here, by orders of magnitude. The federal Black Lung Benefits Program has paid out forty-four billion dollars across half a century — a real number, not a small one — and the program's Trust Fund carries roughly five point one billion dollars in outstanding debt as of September 2024, with no instrument in place that prices the cost severance the trust fund was designed to compensate for. The reclamation bonds posted by mining companies cover four to six billion dollars less than the actual reclamation cost of the sites whose extraction generated the bonds. The intergenerational obligation embedded in the Social Security trust-fund accounting has no Foreclosure Bond against it; the cost is structurally guaranteed to be borne by whichever generation happens to be in working age when the obligation matures. The framework names this as the accountability gap, and the gap is what the next chapters' methodology and policy work address.

The reparations-economics framing carries one further discipline. Reparations economics treats the question of *what is owed* as separate from the question of *what is politically feasible to collect*. The Restitution Bond methodology applies regardless of whether any specific society is currently prepared to post the bond. The framework's contribution, in this regard, is not advocacy — it is accounting. What the bond would have to be, if the cost severance documented in this chapter were accounted for honestly, is the

methodology's question. What political path leads from current bond levels to levels approaching the residual commons value is a separate question, addressed in Chapter 9's policy economy and explicitly scoped, per the framework's discipline, as an open political-philosophical question rather than as a single-instrument prescription this book attempts to legislate.

This chapter has named the cost severance the American extraction regimes produce, documented the empirical pattern across six cases — at financial-system, intergenerational, and institutional-medical scales — engaged the strongest counter-positions a sympathetic reader could press against the structural finding, and introduced the framework's two-instrument response apparatus alongside the academic-and-legal lineages that ground it. What remains is methodology. The next chapter formalizes the apparatus: the Commons Inversion Test that admits cost claims to the framework's accounting; the four gates that filter admitted claims for unit-consistency, boundedness, and independence; the three ways of estimating Residual Commons Value when no single estimation method is sufficient; the Hotelling Identity that bridges the framework to standard resource economics. After methodology, Chapter 7 walks the framework through thought-experiment scaling — does the apparatus hold under maximally-different conditions of off-Earth extraction? — and Chapter 8 walks one ton of Appalachian coal through the framework's full cost-decomposition that produces the chapter's most concrete numerical example. The policy economy that Chapter 9 develops, the closing reflection that Chapter 10 offers, and the technical appendix that operationalizes everything for academic-reviewer engagement all build on the apparatus this chapter has now introduced. The accounting framework exists; what remains is its formalization, its testing, and its honest engagement with the political work of building accountability architecture in societies that have, until now, declined to build it.

That is the question the next three chapters address. Not whether cost severance exists, but what we do about it.

---

— End of Chapter 5 —